

Assignment 1

Submit the GIT repo link as a submission on Openverse

- Given the table:

- orders

id	user_id	status	created_at
----	---------	--------	------------

- Suggest appropriate indexes

- Explain why you chose them

Table

orders

Column	Type (assumed)
id	primary key
user_id	foreign key
status	string / enum
created_at	timestamp

Suggested Indexes

1. Primary Key Index on id

PRIMARY KEY (id)

Why this index is needed

- Every row must be uniquely identifiable.
- Most queries that fetch a single order use the id.
- The database automatically creates this index.

Benefit

- Extremely fast lookup of a single order.
- Ensures no duplicate order IDs exist.

Example

This works like a unique roll number for each order, making it easy and fast to find.

2. Index on user_id

```
CREATE INDEX idx_orders_user_id ON orders(user_id);
```

Why this index is needed

- Applications frequently fetch all orders of a user.
- Common query:

```
SELECT * FROM orders WHERE user_id = ?;
```

- Without an index, the database would scan the entire table.

Benefit

- Faster retrieval of user-specific orders.
- Improves performance when joining with the users table.

Example: Join Query Using user_id

Query

```
SELECT  
    u.id,  
    u.name,  
    o.id AS order_id,  
    o.status,  
    o.created_at  
FROM users u
```

```
JOIN orders o  
ON u.id = o.user_id  
WHERE u.id = 123;
```

Why the `user_id` Index Helps Here

What happens without an index on `orders.user_id`

- Database scans **all rows** in the `orders` table.
- For each order, it checks if `user_id = 123`.
- This is slow when the `orders` table is large.

👉 This is called a **full table scan**.

What happens with an index on `orders.user_id`

```
CREATE INDEX idx_orders_user_id ON orders(user_id);
```

- Database directly jumps to orders with `user_id = 123`.
- No need to scan unrelated orders.
- Join becomes much faster.

👉 This is called an **index lookup**.

Data Size Example

- Users table: 10,000 users
- Orders table: 1,000,000 orders

Without index:

- Database checks all 1,000,000 orders

With index:

- Database checks only the orders belonging to that user (maybe 20–30 rows)
-

3. Composite Index on (user_id, created_at)

```
CREATE INDEX idx_orders_user_created ON orders(user_id, created_at);
```

Why this index is needed

Most applications show a user's orders sorted by date.

Common query:

```
SELECT *
FROM orders
WHERE user_id = ?
ORDER BY created_at DESC;
```

This query does **two things**:

1. Filters rows using user_id
 2. Sorts the result using created_at
-

Case 1: No Index

What the database does:

- Scans the entire orders table
- Filters rows where user_id = 101
- Sorts the filtered result by created_at

Problem:

- Full table scan
 - Expensive sorting step
-

Case 2: Separate Indexes

```
CREATE INDEX idx_user_id ON orders(user_id);
CREATE INDEX idx_created_at ON orders(created_at);
```

What the database does:

- Uses `idx_user_id` to find orders for user 101
- Still needs to **sort the result** using memory or disk
- Cannot efficiently combine both indexes for this query

Problem:

- Sorting still required
 - Extra CPU and memory usage
-

Case 3: Composite Index (Best Case)

```
CREATE INDEX idx_user_created ON orders(user_id, created_at);
```

How the composite index is stored

The data in the index is already ordered like this:

user_id	created_at
101	2024-01-10
101	2024-01-05
101	2024-01-01
102	2024-01-08

What the database does now

- Uses the index to **directly find `user_id = 101`**
- Rows are **already sorted by `created_at`**
- No extra sorting step needed

Result:

- Fast filtering

- Fast sorting
 - Minimal resource usage
-

Simple Analogy

- **Separate indexes:**
Find all pages for “Chapter 3”, then sort them manually.
 - **Composite index:**
Chapter 3 pages are already in order.
-

Performance Benefit Summary

Index Type	Filtering	Sorting	Performance
No index	✗	✗	Very slow
Separate indexes	✓	✗	Medium
Composite index	✓	✓	Fastest